



LUDWIG-
MAXIMILIANS-
UNIVERSITÄT
MÜNCHEN

FAKULTÄT FÜR MATHEMATIK, INFORMATIK UND STATISTIK
INSTITUT FÜR INFORMATIK

**FORSCHUNGSGRUPPE
DATA MINING IN DER MEDIZIN**



Bachelor Thesis
in Computer Science

Attention in Mixed-Type Clustering

Jaanis Fehling

Aufgabensteller: Prof. Dr. Christian Böhm
Betreuer: Walid Durani
Abgabedatum: tt.mm.yyyy

Declaration of Authorship

I hereby declare that the thesis submitted is my own unaided work. All direct or indirect sources used are acknowledged as references.

This paper was not previously presented to another examination board and has not been published.

Munich, tt.mm.yyyy

.....
Jaanis Fehling

Abstract

Clustering mixed-type datasets is a difficult task, and many proposed methods do not manage to beat baseline clustering models that were initially designed for continuous data applications. In addition, they often introduce computational complexity, which makes them impractical in real-world applications, where there is large amounts of data present. With autoencoder neural networks gaining popularity, the possibility of using autoencoders in clustering tasks has emerged. The autoencoder can learn a dense representation of the data, which can then be used for the following clustering step. We extend the autoencoder with an attention mechanism, which allows the autoencoder to learn dependencies between features. Further, we extend the attention mechanism to a full Transformer, which is able to provide contextual feature embeddings. We show that these methods can increase clustering performance on datasets with heavily related features.

Contents

1	Introduction	2
2	Related Work	3
3	Foundation	5
3.1	Methodology	5
3.1.1	Datasets	5
3.1.2	Evaluation	6
3.2	Classical Methods for clustering Mixed-Type data	7
3.2.1	k-means	7
3.2.2	k-modes	8
3.2.3	k-prototypes	9
3.2.4	Gower distance	10
3.3	Deep Clustering	11
3.3.1	Neural Networks	11
3.3.2	Autoencoder	14
3.3.3	Deep Clustering	15
4	Attention in Mixed-Type Clustering	20
4.1	Autoencoder and k-means	20
4.2	Column Embeddings	21
4.3	Attention	22
4.4	Transformer	26
4.4.1	FT-Transformer	26
5	Experiments	30
5.1	Comparison of classical Clustering Methods	30
6	Conclusion	38
	Bibliography	39

Chapter 1

Introduction

Clustering is an unsupervised machine learning method that groups similar observations together. Due to its ability to find patterns in an unlabeled dataset, it is an essential task in Data Mining and Knowledge Discovery. A *cluster* is a group of similar instances that belongs to a *centroid* (center point of a cluster). The goal of this work is to explore clustering of mixed-type datasets, which contain continuous, numeric features as well as categorical features. A wide variety of clustering methods for this domain has been explored [1]. Distance-based clustering algorithms use distance measures such as Euclidean distance to calculate the similarity of data points and assign them to the nearest cluster centroid. Hierarchical methods partition the instances and merge (agglomerative clustering) or split them into bigger or smaller clusters.

Clustering mixed-type datasets is a difficult task, and many proposed methods do not manage to beat baseline clustering models that were initially designed for continuous data applications. In addition, they often introduce computational complexity, which makes them impractical in real-world applications, where there is large amounts of data present.

With autoencoder neural networks gaining popularity, the possibility of using autoencoders in clustering tasks has emerged. The autoencoder can learn a dense representation of the data, which can then be used for the following clustering step. We extend the autoencoder with an attention mechanism, which allows the autoencoder to learn dependencies between features. Further, we extend the attention mechanism to a full Transformer, which is able to provide contextual feature embeddings. We show that these methods can increase clustering performance on datasets with heavily related features.

Chapter 2

Related Work

Lots of research has been put into creating deep neural network classifiers for the tabular data domain that try to beat Gradient boosted decision trees [9]. There has been some research on incorporating attention or a Transformer block into such networks. TabNet [2] is a self-supervised classifier using attention that makes use of feature selection to group features into semantic categories. TabTransformer [18] and FT-Transformer [13] both use a full Transformer block before feeding the tensors into a MLP classifier. TabTransformer uses embeddings only for categorical features and uses a shared embedding to distinguish classes in one column from the classes in other columns. It then uses a Transformer block to attend the categorical features. FT-Transformer also uses embeddings for continuous features but no shared embedding. Using embeddings for both categorical and continuous features allows it to utilize all features in the Transformer block. There has also been research on classifiers that utilize attention not only over the features of one instance, but over multiple instances of a datasets [23, 41].

There has been some research on clustering methods that incorporate an attention mechanism. SADSC [10] leverages convolutional layers combined with self-attention layers for subspace clustering. Methods such as DAGC [36], CMFGN [17] and DLGAMC [25] have used attention in the area of graph clustering.

In the field of deep clustering, research of the last few years has focused on models that are able to learn clustering friendly representations, instead of sequentially connecting an autoencoder and a clustering module. DCN [46] uses a joint loss function to update the autoencoder (representation module), the cluster assignments and centroids. DEC [45] has introduced a whole category of deep clustering methods, which minimize the Kullback-Leibler divergence between the cluster assignment distribution and an auxiliary distribution that is more centered around the cluster centroids. Some methods that fit in this

category are: IDEC [15], which improves DEC by including the reconstruction loss of the autoencoder; DEPICT [11], which uses a convolutional autoencoder; DEC-DA [16], which improves DEC by using data augmentation. Still, DCN and DEC provide the baseline for deep clustering in the image and natural language domain, which newer models are compared against.

Chapter 3

Foundation

3.1 Methodology

3.1.1 Datasets

In this work we use 8 mixed-type datasets from the UC Irvine Machine Learning Repository [29]. The Abalone dataset [32] contains physical measurements from abalones. It has 4177 instances, one categorical feature and seven continuous features. The Auction Verification dataset [33] has 2043 instances that contain verification runs of multi-round auctions. It is composed of six categorical and one continuous feature. The Bank Marketing dataset [31] is related to a direct marketing campaign of a portuguese banking institution. It has 49732 instances, but was downsampled to 5000 random instances. The "age", "day" and "month" features were removed, which results in eight categorical and five continuous features. The Breast Cancer dataset [44] contains 699 instances and 9 categorical features. The Census Income dataset [22] has a total of 48842 instances. It was downsampled to 5000 random instances. It is composed of eight categorical and 6 continuous features. The Credit Approval dataset [38] contains information of applications for credit cards. It has 690 instances, nine categorical features and six continuous features. The Heart Disease dataset [21] is composed of four datasets and has 920 instances in total. It has seven categorical features and six continuous features. The Soybean (Large) Dataset [32] consists of 683 instances from soybeans with a certain disease and has 35 categorical features.

All instances containing missing values were removed. Duplicate instances were explicitly not removed, since there is no information available if they are duplicates by accident, or real duplicates. Categorical features were standardized by removing the mean and scaling to unit variance, using the scikit-learn Python library [35]. Formally, the standardized score z of a sample x from a

feature is calculated as

$$z = \frac{(x - \mu)}{\sigma},$$

where μ is the mean of the samples x_1, \dots, x_N from a feature of length N , defined as

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i$$

and σ is the standard deviation of the samples of a feature, defined as

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}.$$

The datasets were shuffled. For all random operations, a random state of integer value 0 was used to ensure reproducibility. When using the PyTorch Python library [34] for implementing neural networks, a flag was set to only use deterministic algorithms.

3.1.2 Evaluation

The clustering results were each evaluated with two measurements, *accuracy* and *Normalized Mutual Information*. Accuracy is defined as

$$\frac{\text{Instances predicted correctly}}{\text{Total number of instances}}.$$

Since any clustering algorithm will assign arbitrary class labels to each instance that might not align with the ground truth class labels, we map the predicted label classes to ground truth label classes in a way maximize the number of correctly predicted instances. This is an inversion of the *Linear Assignment Problem*, we use an implementation provided by the Scipy Python library [43] to solve this problem and find the optimal solution.

As shown in Chapter 5, some clustering methods falsely assign almost all instances to one target class. Because a part of the datasets we used are heavily imbalanced, this leads to unjustified high accuracy scores. Therefore, we use another metric, Normalized Mutual Information. It is based on *Mutual Information*, which, for the ground truth labels U and the predicted labels V (switching both variables will not change the outcome) is defined as

$$\text{MI}(U, V) = \sum_{i=1}^{|U|} \sum_{j=1}^{|V|} \frac{|U_i \cap V_j|}{N} \log \frac{N|U_i \cap V_j|}{|U_i||V_j|}$$

where $|U_i|$ is the number of instances assigned to cluster U_i and $|V_i|$ is the number of instances assigned to cluster V_i [35]. Normalized Mutual Information is Mutual Information normalized by the mean of the Entropies of U and Y [35]:

$$\text{NMI}(U, V) = \frac{\text{MI}(U, V)}{\frac{1}{2}(H(U)H(V))}.$$

3.2 Classical Methods for clustering Mixed-Type data

3.2.1 k-means

The most well known distance-based clustering method is k-means [27]. The goal is defined as follows: Suppose we have a finite set of n instances $S = \{p_1, p_2, \dots, p_n\} \in \mathbb{R}^m$ for a dataset with m features, the target of k-means is to find optimal centroids $B = \{b_1, b_2, \dots, b_k\} \subseteq \mathbb{R}^m$ for a given $k(\leq n) \in \mathbb{N}$ that minimize the sum of the squared Euclidean distance of each point in S to its nearest centroid. Formally

$$\sum_{i=1}^n d(p_i, B)^2$$

has to be minimized, where d is the Euclidean distance from a point $p_i \in S$ to the nearest centroid in B [28]:

$$d(p_i, B) = \min_{1 \leq j \leq k} d(p_i, b_j).$$

The Euclidean distance between two points p and q in an n -dimensional Euclidean space is defined as

$$d(p, q) = \|p - q\| = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_n - q_n)^2}.$$

Finding the optimal centroids is a NP-hard problem, even for $d = 2$, as shown by Mahajan et al. [28]. The most common algorithm used for the k-means problem is a iterative refinement technique proposed by Lloyd [26]. It is defined as follows

1. Randomly set k initial cluster centroids $b_1^{(1)}, \dots, b_k^{(1)}$.
2. Assign each observation p_i to the nearest centroid using squared Euclidean distance. This splits our instances into S into k sets $\{S_1^{(t)}, \dots, S_k^{(t)}\}$.
3. Recalculate the optimal position of each centroid using the mean distance to each instance assigned to the centroid:

$$b_i^{(t+1)} = \frac{1}{|S_i^{(t)}|} \sum_{p_j \in S_i^{(t)}} p_j.$$

4. Repeat steps 2. and 3. until the centroid assignments no longer change.

3.2.2 k-modes

In many real-world scenarios, besides continuous, numerical data, *categorical* data exists. While Euclidean distance or other distance measures work well with continuous data, categorical data is different. Suppose we have categories $\{A, B, C\}$ of a given feature, we would encode them into numeric values to allow for computation of a distance measure:

$$\{A, B, C\} \equiv \{1, 2, 3\}.$$

While A and C can share the same semantic similarity as A and B , numerically category A and Category C are now $|1 - 3| = 2$ apart, while Category A and B are only $|1 - 2| = 1$ apart. During clustering, this could lead to instances being assigned to centroids based on a wrong distance assumption.

A possible solution is to use *one-hot encoding*, also known as *dummy coding* in classical statistics. One-hot encoding turns a discrete feature containing k mutually exclusive categories into a vector x of length k , in which only one of the elements x_k equals 1 and all remaining elements equal 0 [7]. For an instance B of a feature having $k = 3$ separate categories $\{A, B, C\}$, the one-hot vector x would be represented by $x = (0, 1, 0)^\top$. According to Huang [19], one-hot encoding has two drawbacks:

1. In real-world applications, categorical features with hundreds or thousands of categories are encountered. This would result in a large number of binary features in the one-hot encoded representation, which will increase cost and space of computation.
2. The centroid value of a certain one-hot encoded feature, given by a real value between 0 and 1, cannot indicate the characteristics of the according cluster, since the feature only describes the presence or absence of one category.

Therefore, Huang [19] proposed using the Kronecker-Delta as a dissimilarity measure between multiple categorical features. Formally, if we have two instances X and Y of a dataset with m categorical features, d_1 will count the number of mismatches between the categorical features of both instances, defined as

$$d_1(X, Y) = \sum_{j=1}^m \delta(x_j, y_j)$$

where the Kronecker delta $\delta(x_j, y_j)$ is defined as

$$\delta(x_j, y_j) = \begin{cases} 0 & (x_j = y_j) \\ 1 & (x_j \neq y_j) \end{cases}.$$

If we have a finite set of n instances $S = \{p_1, p_2, \dots, p_n\}$ for a dataset with m categorical features, the goal of k -modes [19] is to find optimal *modes* $B = \{b_1, b_2, \dots, b_k\}$ for a given $k(\leq n) \in \mathbb{N}$ that minimize

$$\sum_{i=1}^n d_1(p_i, B)$$

where

$$d_1(p_i, B) = \min_{1 \leq l \leq k} \sum_{j=1}^m \delta(p_{i,j}, b_{l,j}).$$

Similar to k-means, we can use an iterative algorithm for efficient computation [19]:

1. Randomly choose k instances from the dataset as initial modes for the clusters.
2. Assign each instance to their nearest mode using the proposed dissimilarity measure one by one and update the mode of each cluster after each assignment.
3. Test if each instance still belongs to its assigned mode, i.e. if each instance is assigned to its nearest mode. If the instance would belong to a different mode, reassign the instance and update the modes of both clusters.
4. Repeat step 3. until the mode assignments no longer change.

3.2.3 k-prototypes

As proposed by Huang [19], it is straightforward to combine the k-means and k-modes algorithms into the *k-prototypes* algorithm, which can be used to cluster *mixed-type* data (consisting of numerical, continuous and categorical features). The dissimilarity between two instances X and Y with features $A_1^r, A_2^r, \dots, A_s^r, A_{s+1}^c, \dots, A_m^c$, where features A_1^r, \dots, A_s^r are continuous and features A_{s+1}^c, \dots, A_m^c are categorical, is defined as

$$d_2(X, Y) = \sum_{j=1}^s (x_j - y_j)^2 + \gamma \sum_{j=s+1}^m \delta(x_j, y_j).$$

The first part of the equation is the Euclidean distance as used in k-means, while the second part is taken from the k-modes algorithm. Huang [19] states: "The weight γ is used to avoid favouring either type of attribute".

Again, we need to find k optimal centroids $B = \{b_1, b_2, \dots, b_k\}$ and therefore have to minimize

$$\sum_{i=1}^n d_2(p_i, B)$$

where

$$d_2(p_i, B) = \min_{1 \leq l \leq k} \sum_{j=1}^s (p_{i,j} - b_{l,j})^2 + \gamma \sum_{j=s+1}^m \delta(p_{i,j}, b_{l,j}).$$

We can minimize both distance measures at the same time since they are nonnegative. Therefore, we can use the same algorithm as defined in 3.2.2. [19]

3.2.4 Gower distance

Gower distance [14] is a general similarity measurement between instances containing mixed-type features. It is defined as follows: When comparing instances x_i and x_j , for each feature k of p total features, we calculate a score $s_{ijk} \in [0, 1]$. The score will be close to 1 for two instances x_{ik} and x_{jk} of a feature k if they are similar, and close to 0 they are not similar. Gower distance is also computable between instances with missing values, therefore a quantity δ_{ijk} is calculated, which is equal to 1, when feature k can be compared across the two instances x_i and x_j , and 0 otherwise (illustrated in Figure 3.1). Gower distance then is the average of the known score

$$S_{ij} = \frac{\sum_{k=1}^p s_{ijk} \delta_{ijk}}{\sum_{k=1}^p \delta_{ijk}}.$$

The Score s_{ijk} is calculated differently according to the type of feature [14]:

1. For *dichotomous* (when a value is either present or absent) features, the score s_{ijk} is 1 when the value is present in both features and 0 otherwise, as shown in Figure 3.1.
2. For categorical features, the score s_{ijk} is 1 if they both instances match on feature k and 0 otherwise.
3. For continuous features the score is calculated as

$$s_{ijk} = 1 - \frac{|x_i - x_j|}{R_k}$$

where R_k is the range of feature k in the dataset or in the sample.

SCORES AND VALIDITY OF DICHOTOMOUS CHARACTER COMPARISONS

Individual i j	Values of character k			
	+	+	-	-
	+	-	+	-
s_{ijk}	1	0	0	0
δ_{ijk}	1	1	1	0

Figure 3.1: Illustration from the original paper by Gower [14]. Score s_{ijk} and quantity δ_{ijk} of a feature k on two instances x_i and x_j . Presence of a feature is denoted by "+" and absence by "-".

Gower [14] has shown that $\sqrt{1 - S_{ij}}$ is a valid distance representation for two instances x_i and x_j . We can now convert our similarity matrix S into a distance matrix and are able to use Hierarchical clustering methods [20]. Philip and Ottaway [37] used Gower distance with agglomerative clustering. Agglomerative clustering places each instance into its own cluster and recursively merges the clusters together using the given distance matrix, until only the specified number of clusters is remaining [20].

3.3 Deep Clustering

3.3.1 Neural Networks

The idea of computation by neurons inspired by the human brain was first formalized into a mathematical model by McCulloch [30]. A *neuron* was defined as a element that takes multiple boolean inputs and has one boolean output. The neuron *fires*, meaning the output is set to true, when the sum of the input values extends a certain threshold.

The single-layer *perceptron*, the first neural machine learning algorithm, was invented by Rosenblatt in 1957 [39]. Formally, the binary valued output o_j given an input vector x_i is calculated as

$$o_j = \begin{cases} 1 & \sum_i w_{ij}x_i + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

where w is the learnable weight matrix, and b is a predefined bias. For training, the weights w are simply incremented when the output is 0 but the ground truth is 1, and decremented if the output is 1 but should be 0. If the output was predicted right, no weights are changed.

The idea of neural networks was revived three decades later, using multiple perceptron layers and a differentiable error function [40]. In a multi-layer neural network, we have an input layer, an output layer and multiple hidden layers. Each layer is a collection of neurons, that acquire the outputs of each neuron from the previous layer (or from the input in case of the input layer) as their input, and produce a new output. Formally, the output a_j of the j th neuron of layer n that gets d input values from the previous layer is defined as

$$a_j = \sum_{i=1}^d w_{ji}^{(n)} x_i + b_j^{(n)}$$

where $w_{ji}^{(n)}$ is the learnable weight of the input x_i going through neuron j in layer n [6]. A bias b_j is also added. This output is commonly referred to as an *activation* of a neuron [6]. Because every neuron is a linear function, in order to avoid the collapse of each neuron from all layers into a single linear function, we pass the activation a_j into a non-linear activation function f

$$z_j = f(a_j)$$

before being passed to the next layer of neurons [12].

As shown in a recent survey [12], there are many different activation functions used in neural networks. A simple example for an activation function, which naturally aligns with the idea of a biological neuron firing is the *step function* (also known as *Heaviside step function*) [6], that outputs 0 for negative values and 1 for positive values:

$$\text{step function}(x) = \begin{cases} 1 & x \geq 0 \\ 0 & x < 0 \end{cases}.$$

In order to stay between 0 and 1, but also utilize all real values in between, another activation function that has historically been popular is the *logistic sigmoid* function [12], that squashes any input in $]0, 1[$:

$$\text{logistic sigmoid}(x) = \frac{1}{1 + e^{-x}}.$$

Because the outputs are in a range close to 0, it is prone to the *vanishing gradient problem*. In the vanishing gradient problem, a gradient that is very




Figure 3.2: An illustration of the step function, the logistic sigmoid, the ReLU and the Leaky ReLU activation functions.

close to 0 leads to almost no update in the weights of the network during training. Moreover, using an exponential value naturally leads to a greater computational complexity. [12]

The current state-of-the-art activation function, the *Rectified Linear Unit* (ReLU) [12], is not limited by the above disadvantages. It is simple and computationally performant. It is defined as the identity for positive values and as 0 for negative values:

$$\text{ReLU}(x) = \max(0, x) = \begin{cases} x & x \geq 0 \\ 0 & x < 0 \end{cases}.$$

One downside of the ReLU activation function is that there will be a substantial amount of dead neurons in the network (neurons with output 0 will not affect the neurons of the next layer). The *Leaky Rectified Linear Unit* (Leaky ReLU or LReLU) [12] utilizes negative values as well, but with a small coefficient, usually set to 0.01. It is defined as

$$\text{Leaky ReLU}(x) = \begin{cases} x & x \geq 0 \\ 0.01 \times x & x < 0 \end{cases}.$$

All four activation functions are illustrated in Figure 3.2.

In a multi layer neural network, we need to be able to update the weights during training, so the network can learn the most optimal weights in order to fulfill its designated task. In the single layer perceptron, we were able to change the weights directly based on the predicted output, but in a multi layer network this is non-trivial, since we have multiple such perceptron layers that are sequentially attached. The solution is to use differentiable activation functions (or activation functions differentiable at almost every point, e.g. ReLU is not differentiable at 0), so each neuron together with its activation function becomes a differential function of the input, the weight and bias. We can now define a differentiable error function of the network outputs, which therefore is also a differentiable function of the weights. A common error function used in neural networks is the *mean squared error*, defined as

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - z_i)^2,$$

where n is the number of input instances, y_i is the target label of instance x_i and z_i is the network output on instance x_i . The derivatives of the error function in combination with stochastic gradient descent can be then be used to find the weights that minimize the error function. [6]

3.3.2 Autoencoder

A neural network, as defined in the previous section, needs the target labels of a dataset in order be trainable. In a unsupervised learning setting, this is unapplicable. An approach used for pre-training neural networks with unlabeled data is using an *autoencoder* neural network, first formalized by Ballard [5]. An autoencoder is a neural network, that projects the input into a lower-dimensional embedding space, and then back into its original dimensionality. The network tries to reconstruct the original input the best it can, after it was encoded into the embedding space. Therefore the target label fed into the loss function is the original input. This allows such a network to be trained without the need for labeled data. The autoencoder consists of two blocks, where each is a collection of fully connected linear layers. The *encoder* maps the input into the embedding space, while the *decoder* projects the encoded representation into its original dimensionality. Both blocks can contain multiple hidden layers, but typically they are constructed symmetrical, meaning the number and size of the linear layers in the decoder are the same as in the encoder, but inverted (illustrated in Figure 3.3).

In unsupervised learning such as clustering, the autoencoder can be used to learn a dense representation of the input data. After training, we can feed the data into the encoder part of the autoencoder in order to obtain the

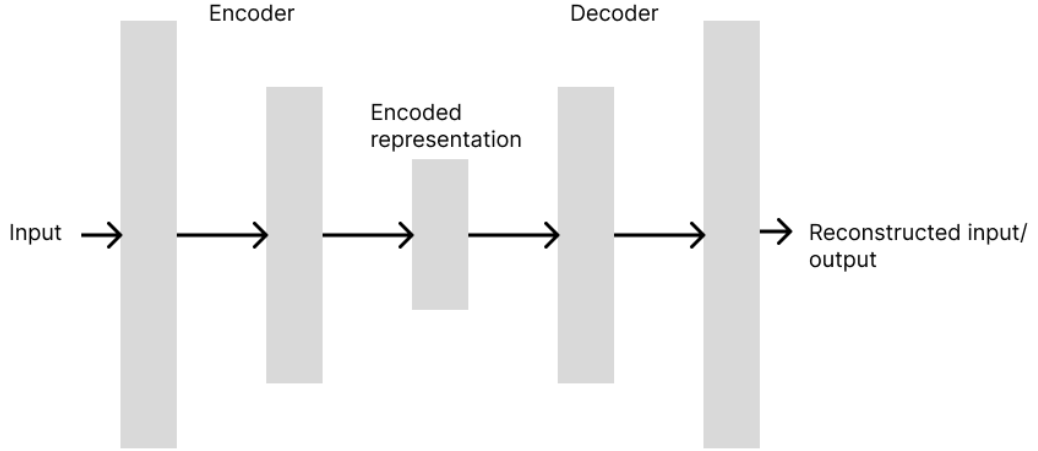


Figure 3.3: Illustration of an autoencoder neural network.

encoded data, before utilizing a classical clustering algorithm to acquire the final cluster labels. The encoded representation has a smaller dimensionality and only contains the most relevant parts of the data (which were needed for reconstruction), which can benefit clustering.

3.3.3 Deep Clustering

The problem of training an autoencoder before clustering is that the autoencoder cannot learn a suitable representation for the subsequent clustering step. The autoencoder has no objective that rewards learning a cluster-friendly representation [46]. Instead of using a multi-stage process that splits autoencoder training and clustering, recent research has focused on a variety of solutions that combine feature representation learning and clustering [47]. The general goal is to learn feature representations that are suitable for clustering.

Xie et al. [45] proposed DEC (Deep Embedded Clustering), an approach using an iterative process that simultaneously learns feature representations and cluster assignments. DEC initially trains the autoencoder on the data as with the previous multi-stage approach, but then leaves out the decoder, as illustrated in Figure 3.4, and trains the encoder separately while clustering to learn more optimal feature representations. After the initial autoencoder pre-training, standard k-means (Section 3.2.1) is used on the encoded representation to obtain initial cluster labels. A two-step algorithm is then used





Figure 3.4: Illustration of DEC by Xie et al. [45].


 Figure 3.5: Illustration from Li [24] on the distributions Q and P on two cluster centroids μ_1 and μ_2 . It is notable that P is more concentrated around the cluster centers, but still derived from Q .

to iteratively update both the cluster assignments and feature representations. The first step involves computing a *soft assignment* between the encoded data points and the cluster centroids. Student's t -distribution is used as a kernel, with fixed degrees of freedom $\alpha = 1$:

$$q_{ij} = \frac{(1 + \|z_i - \mu_j\|^2/\alpha)^{-\frac{\alpha+1}{2}}}{\sum_{j'} (1 + \|z_i - \mu_{j'}\|^2/\alpha)^{-\frac{\alpha+1}{2}}},$$

where z_i represents the encoded value of instance x_i and μ_j the centroid value of cluster j . The resulting value q_{ij} can be interpreted as the probability of assigning instance x_i to cluster j . The objective of the second step is minimizing the Kullback-Leibler divergence between our previously defined distribution Q and an auxiliary distribution P :

$$L = \text{KL}(P\|Q) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}.$$

The auxiliary distribution P is more concentrated around the cluster centers than Q (illustrated in Figure 3.5), but not using an unnatural distribution such as a delta distribution. The goal of P is to strengthen clustering predictions, put more emphasis on points assigned with high confidence and normalize loss contribution of each centroid. Therefore, P is defined as

$$p_{ij} = \frac{q_{ij}^2/f_j}{\sum_{j'} q_{ij'}^2/f_{j'}},$$

where $f_j = \sum_i q_{ij}$ is the soft cluster frequency. We can then train the encoder using stochastic gradient descent to produce points z_i that minimize the Kullback-Leibler divergence. [45]

Guo et al. [15] proposed IDEC (Improved Deep Embedded Clustering), which tries to improve DEC by also using the autoencoder reconstruction loss. Formally, the autoencoder reconstruction loss L_r is added to the DEC loss $L_c = \text{KL}(P\|Q)$, which gives:

$$L = L_r + \gamma L_c,$$

where $\gamma > 0$ is a hyperparameter. This is illustrated in 3.6.

Yang et al. [46] proposed DCN (Deep Clustering Network) (Figure 3.7), which uses a joint loss function to optimize both feature representations and clustering predictions. The loss function is defined as

$$L = \sum_{i=1}^N \left(\ell(g(f(x_i)), x_i) + \frac{\lambda}{2} \|f(x_i) - Ms_i\|_2^2 \right).$$




Figure 3.6: Illustration of IDEC by Guo et al. [15].




Figure 3.7: Illustration of DCN by Yang et al. [46].

The loss function consists of two parts, the reconstruction loss (first term of the sum) and the clustering loss (second term of the sum). λ is used as a hyperparameter to balance the two terms. The network then alternates between updating the network parameters, updating the assignments to the cluster centroids and updating the cluster centroids themselves.

Chapter 4

Attention in Mixed-Type Clustering

4.1 Autoencoder and k-means

It is difficult to formulate statistical distance measures for comparing two values of a categorical feature. Categorical features are irregular in nature, they can be binary, ordinal, non-ordinal or composed of many semantically different classes. This makes it hard to cluster a mixed-type datasets with classical clustering methods. As described by Huang [19] and also explained in Section 3.2.2, one-hot encoding is suboptimal for the use in clustering tasks. The proposed solution by Huang, k-prototypes [19], does not beat k-means or k-means with one-hot encoding on any tested dataset, as shown in Chapter 5. Gower distance shows great accuracy results on all datasets and manages to beat the other methods on five out of eight datasets. As further illustrated in Chapter 5, the accuracy evaluation metric is highly problematic, since some datasets are heavily imbalanced. In fact, the datasets on which gower distance achieves the best accuracy results are greatly imbalanced. This means the agglomerative clustering used with gower distance falsely assigns almost all instances to one target class. When looking at Normalized Mutual Information (NMI), gower distance with agglomerative clustering does not beat k-means or k-means with one-hot encoding on any dataset. This leaves us with naive k-means intended for only numerical, continuous data and k-means with one-hot encoding, which is a inefficient, suboptimal encoding.

Therefore, instead of using statistical distance measure to compare instances with categorical and continuous features, we use an autoencoder neural network to learn a dense representation of the data. The clustering architecture we build upon is a multi-stage process. We first train the autoencoder on our dataset independently of the clustering process. The training step is repeating

many times, in our experiments we train the network a hundred times (hundred epochs). After the initial training step, we can feed the dataset into the encoder part of our trained autoencoder to get the dense, encoded representation of our data. We can perform k-means clustering over this representation to get the final cluster labels.

4.2 Column Embeddings

Even when using neural networks, there is still the problem of representing categorical values, especially non-ordinal values, as explained in depth in Section 3.2.2. Therefore, we use a learnable *embedding* layer for categorical values. The resulting embedded tensor of the embedding layer for a categorical feature is a meaningful representation that is more dense than a one-hot encoded feature. The embedding layer is composed of a weight tensor e_i for each categorical feature i . The weight tensor e_i has the following dimensionality: $|i| \times s_i$, where $|i|$ is the number of unique classes in categorical feature i and s_i is the corresponding embedding dimension. We calculate the embedding dimension s_i of feature i as

$$s_i = \min(50, \lceil \frac{|i|}{2} \rceil).$$

We then encode the classes of each categorical feature with integer values, going from 0 onwards. These integer values serve as indices for the corresponding weight tensor of our embedding layer. Passing a value j of a feature i in the embedding layer returns the one-dimensional tensor at index j of the weight tensor e_j , which is of dimension $1 \times s_i$. The resulting tensors of each categorical feature from the embedding layer are concatenated into a tensor of size $1 \times \sum_{i=1}^n s_i$, where n is the number of categorical features. To construct the complete input tensor, we concatenate the tensor of the embedded categorical features with the tensor of the continuous features. The continuous features of an instance are not embedded, but scaled by removing the mean and scaling to unit variance as explained in Section 3.1.1. Having m continuous features would result in a tensor of continuous features with size $1 \times m$. After concatenating both tensors, this leaves us with an input tensor of size $1 \times (m + \sum_{i=1}^n s_i)$. This input tensor can then be passed into the first linear layer of the encoder.

The goal of the autoencoder therefore is to reconstruct the embeddings, not the actual input. To be able to use a loss function on this network, we clone the tensor of the embedded categorical features and save it in a variable, and later compare the decoder output with the cloned embedded tensor (concatenated with the continuous features).

4.3 Attention

Self-attention is a machine learning mechanism for sequence processing, originally used in machine translation by Bahdanau et al. [4] and extended into the *Transformer* architecture by Vaswani et al. [42]. It is a sequence-to-sequence operation, that aims to create a better representation of the input sequence by putting each element into the context of the other elements (paying "attention" to each other element). We do not operate with a sequence, but instead with multiple categorical features that might be related. For example, the relevance of a feature for the clustering process might depend of the value of another feature. The goal of the attention mechanism is to augment this feature (increase or decrease its vector representation) only when the other feature has this certain value.

Self-attention operates with three input tensors: *queries*, *keys* and *values*. Each of them is the input tensor, but used for different parts of the attention process. We add three independent linear projection to the input tensor to learn different representations for queries, keys and values. We apply attention to the categorical features after they were embedded in the embedding layer and then projected into queries, keys and values. The resulting vectors for each categorical feature are called *contextual embeddings*, since they contain information on the context of the input sample. We concatenate these embedding vectors lengthwise, before passing them into the first linear layer of the encoder, as illustrated in Figure 4.1.

Queries, keys and values are a sequence of n vectors that have d dimensions, resulting in dimensionality $n \times d$. In our case, the input dimension and therefore also the queries Q , keys K and values V are one-dimensional vectors, since the embedded vectors of each categorical features are concatenated along the first axis, as further illustrated in the previous section. We therefore cannot use our previously created input tensor. Because we calculate an individual embedding dimension (with a formula defined in the previous section) for each categorical feature, we cannot concatenate the vectors row-wise, since they all differ in length. Instead, we have to use a fixed embedding dimension d for every categorical feature. Then, we can create tensors Q , K and V with dimension $n \times d$ each, where n is the number of categorical features. The value 32 was chosen for embedding dimension d .

Because Attention is a mechanism that calculates values that combine multiple features, we need to be able to differentiate between the classes of different features, since we have a separate learnable embedding for each feature. A solution in the domain of text translation proposed by Vaswani et al. [42] is to add a positional encoding to each embedding in order to indicate the position of the word in the sentence. This is not suitable for tabular data, since the




Figure 4.1: Architecture of the self-attention autoencoder network.

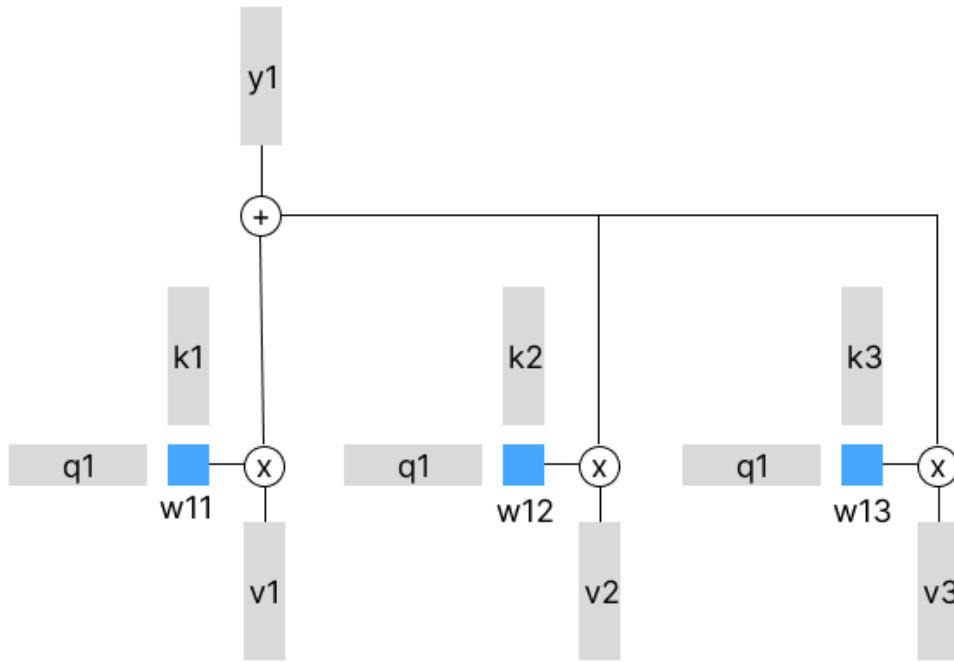


Figure 4.2: Illustration of the attention mechanism on calculating the first output vector, inspired by Bloem [8]. q_1 , k_1 , v_1 represent the vectors at index 1 of Queries Q , Keys K and Values V respectively. The softmax function over the weights is not part of this illustration.

features do not have a particular order. We instead use a solution proposed by Huang et al. [18], which uses a shared embedding across all categorical features. The goal of the shared embedding is to learn dissimilar representations for different classes of separate features in order to distinguish them. The shared embedding is a learnable tensor of size $n \times \frac{d}{8}$, that is concatenated to our input tensor in the first dimension. The fixed embedding dimension d is therefore subtracted by 8, which results in a input tensor of categorical features of size $n \times d$.

The first step of self-attention is to calculate attention weights for every element of the input sequence. For that, we will multiply every query of our queries with the keys. Formally, we can define this as a matrix multiplication:

$$W = QK^\top,$$

where W are the resulting attention weights, Q are the queries and K the keys. The goal of this operation is to calculate weights W_{ij} that indicate how important feature i is to feature j and vice versa. Each weight is an arbitrary value, therefore we use the softmax function [6] to flatten the weights. The softmax function on a weight w_{ij} is defined as

$$\text{softmax}(w_{ij}) = \frac{\exp(w_{ij})}{\sum_j w_{ij}}.$$

Because the weights W can grow towards positive and negative infinity, in order to minimize the vanishing gradient problem of the softmax function (further explained in Section 3.3.1 in the case of the sigmoid function), we scale the attention weights by \sqrt{d} before passing them into the softmax function:

$$W = \frac{QK^\top}{\sqrt{d}}.$$

We are scaling by \sqrt{d} , since its the average increase in Euclidean length of a weight vector w when increasing its dimension d . [42]

Now after we have obtained the attention weights that indicate how important a feature is to every other feature, we need to multiply these weights with each feature vector. This will result in vectors that carry information to which other feature vectors we need to pay attention to. We use the remaining linear projected tensor, the values V , for this step. The full Attention mechanism is therefore defined as

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V.$$

An illustration on how the first output vector in the attention mechanism is calculated is given by Figure 4.2.

4.4 Transformer

The *Transformer* is a architecture based on the attention mechanism, first formalized by Vaswani et al. [42]. The Transformer extends attention into *multi-head attention* and combines it with a feed-forward network and layer normalization. Multi-head attention [42] splits the input in h parts and learns h different representations for queries Q , keys K and values V respectively. Attention is then performed h times in parallel using the different representations of Q , K and V . This allows the model to perform attention over different subspaces of the embedded features at the same time. Using normal attention, we would only be able to attend to the complete embedding of the other features.

After performing attention, the Transformer adds the attention outputs to the initial input and normalizes this sum using *layer normalization* [3]. The output is then fed into two linear layers, the first layer projecting the input to four times its size, and the second layer projecting it onto the original size again. Once again, the output is added to the input of the linear layers, and normalized using layer normalization. This process, illustrated in Figure 4.3, is sequentially applied N times. The layer normalization and feed-forward network between each multi-head attention operation allow the model to parameterize the attention outputs. This makes sure the model can learn from each attention step.

The Transformer was originally designed for natural language tasks, but Huang et al. [18] proposed the *Tab Transformer*, a multi-layer perceptron classifier for the tabular domain, based on the Transformer architecture. Inspired by the performance increase on mixed-type datasets using this architecture, we build a Transformer-based autoencoder network for clustering. The architecture, illustrated in Figure 4.3, is similar to the previous attention-based autoencoder but uses $N = 6$ Transformer blocks instead of one single self-attention block. Our categorical feature embedding size is 32 as previously. We perform multi-head attention with $h = 8$ heads, yielding a embedding dimension for each head of $\frac{d}{h} = 4$.

4.4.1 FT-Transformer

Gorishniy et al. extends the Tab Transformer into the *FT-Transformer*, which uses embeddings for the numerical features as well. We implement a FT-Transformer autoencoder network, illustrated in Figure 4.4. Our implementation uses separate linear layers for each continuous column, that transform a singular value of a continuous feature into a $1 \times d$ vector representation, where d is the embedding dimension of the categorical features. This allows the contin-




Figure 4.3: Architecture of the Transformer autoencoder network.

uous embeddings to be concatenated with the categorical embeddings, before being passed into the Transformer block. The rest of the implementation remains canonical to the previous architecture, only we do not concatenate the continuous feature values with the Transformer output, since the continuous features are already part of the Transformer block. The FT-Transformer paper does not use a shared embedding, we decided to keep still implement it, since Huang et al. [18] showed great accuracy gains when using a shared embedding.

The intuition behind using the FT-Transformer is that many continuous features in the tabular domain are only composed of a small range of values, and can therefore easily represented by an embedding vector. Using Attention over both continuous and categorical features allows the model to learn dependencies across both types of features.




Figure 4.4: Architecture of the FT-Transformer autoencoder network.

Chapter 5

Experiments

5.1 Comparison of classical Clustering Methods

Our first comparison is between classical clustering methods for clustering mixed-type data. We evaluated k-means, k-means with one-hot encoding, k-prototypes and gower distance with agglomerative clustering, using average linkage (average distance between each instance of two sets). The results are given by Figure 5.1 and 5.2. k-means without any modification for mixed type data shows great Normalized Mutual Information (NMI) scores, k-means with one-hot encoding can improve these on some datasets, but loses on others. k-prototypes can beat k-means only on the Census Income dataset, which does not make it a valid choice for clustering mixed-type data. Gower distance, as explained in Section 3.1.2, falsely assigns almost all data points to one cluster centroid, which due to the high class imbalance in our datasets as shown in

	k-means	k-means one-hot	k-prototypes	Gower distance
Abalone	0.1718	0.1740	0.1716	0.1614
Auction Verification	0.0162	0.0071	0.0077	0.0062
Bank Marketing	0.0198	0.0261	0.0195	0.0013
Breast Cancer	0.7468	0.7363	0.5925	0.5537
Census Income	0.1080	0.1850	0.1417	0.0043
Credit Approval	0.3131	0.1710	0.1166	0.0035
Heart Disease	0.2046	0.1645	0.1893	0.1408
Soybean Disease	0.6722	0.7102	0.5676	0.6695

Figure 5.1: Comparison of Normalized Mutual Information of various classical methods on clustering mixed-type datasets.

	k-means	k-means one-hot	k-prototypes	Gower distance
Abalone	0.1353	0.1314	0.1343	0.1954
Auction Verification	0.6647	0.5761	0.5805	0.8008
Bank Marketing	0.7796	0.7866	0.7872	0.8842
Breast Cancer	0.9605	0.9502	0.9151	0.9004
Census Income	0.6082	0.6976	0.6256	0.7684
Credit Approval	0.8086	0.7060	0.6662	0.5482
Heart Disease	0.3344	0.3211	0.4247	0.5652
Soybean Disease	0.5765	0.5996	0.4715	0.501

Figure 5.2: Comparison of accuracy of various classical methods on clustering mixed-type datasets.

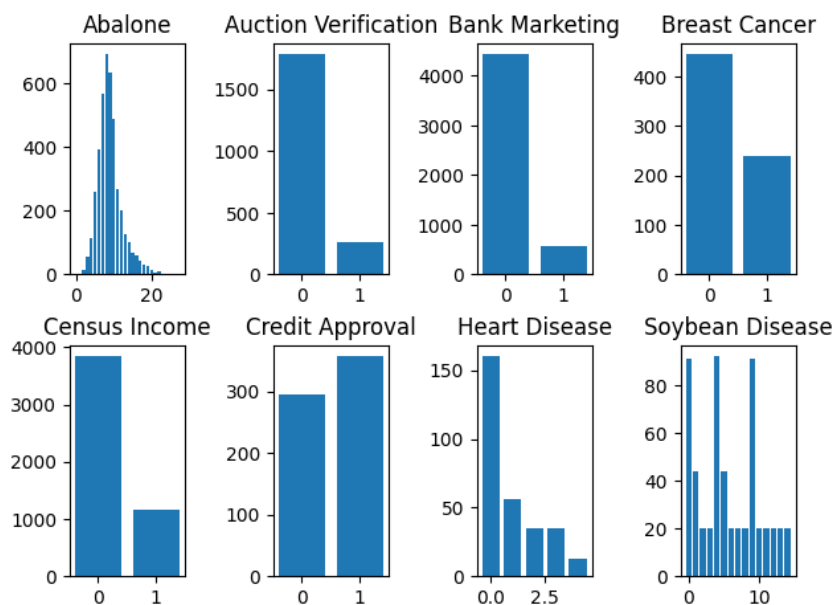


Figure 5.3: Number of instances of each target class of each dataset used.

	No categorical embedding AE	All features AE	Categorical features AE	All features to categorical features AE
Abalone	0.1559	0.1695	0.1684	0.1709
Auction Verification	0.0066	0.0012	0.0002	0.0413
Bank Marketing	0.0015	0.0039	0.0197	0.0000
Breast Cancer	0.7602	0.7178	0.2878	0.6927
Census Income	0.0003	0.0241	0.1024	0.0042
Credit Approval	0.0031	0.0097	0.0008	0.0028
Heart Disease	0.1389	0.1663	0.1848	0.0947
Soybean Disease	0.4974	0.4836	0.4169	0.5426

Figure 5.4: Comparison of Normalized Mutual Information of various autoencoder architectures combined with k-means on clustering mixed-type datasets.

	No categorical embedding AE	All features AE	Categorical features AE	All features to categorical features AE
Abalone	0.1154	0.1688	0.1384	0.1348
Auction Verification	0.5497	0.5893	0.6667	0.8243
Bank Marketing	0.5310	0.7420	0.7894	0.7484
Breast Cancer	0.9634	0.9502	0.7013	0.9458
Census Income	0.7328	0.6664	0.5778	0.6216
Credit Approval	0.5069	0.5176	0.5130	0.5391
Heart Disease	0.3244	0.3278	0.4013	0.3545
Soybean Disease	0.3986	0.4075	0.3292	0.4413

Figure 5.5: Comparison of accuracy of various autoencoder architectures combined with k-means on clustering mixed-type datasets.

Figure 5.3, yields high accuracy scores. Therefore more attention should be paid to NMI Scores, on which gower distance does not beat k-means or k-means with one-hot encoding on any dataset.

The next comparison (Figure 5.4 and 5.5) revolves around the autoencoder architecture that will be chosen for following modifications. We have the choice between only learning dense representations for categorical features or learning them for all features. Further, we can leave out the categorical embeddings and feed the categorical values as they are into the autoencoder. An autoencoder that gets all features as input in the encoder but only learns to reconstruct the categorical features is also added to the comparison. In order to equalize the autoencoders across all datasets and architectures, the hidden encoded dimension d_h of the autoencoder is calculated by

$$d_h = \max(1, \min(\lfloor d_x/4 \rfloor, \lfloor d_y/4 \rfloor)),$$

	AE with categorical embeddings	AE with Attention	Transformer N=6	FT-Transformer N=6
Abalone	0.1603	0.1681	0.1612	0.1750
Auction Verification	0.1234	0.0109	0.0045	0.0065
Bank Marketing	0.0040	0.0394	0.0022	0.0040
Breast Cancer	0.5521	0.1906	0.5569	0.4371
Census Income	0.1475	0.0237	0.0056	0.1423
Credit Approval	0.0123	0.1833	0.3061	0.0897
Heart Disease	0.1550	0.1244	0.1863	0.1717
Soybean Disease	0.5656	0.2514	0.3966	0.2755

Figure 5.6: Comparison of Normalized Mutual Information between an autoencoder with categorical embeddings, an autoencoder using the attention mechanism, a 6-layer Transformer with an autoencoder and a 6-layer FT-Transformer with an autoencoder on clustering mixed-type datasets using k-means.

where d_x is the input dimension of the encoder and d_y the output dimension of the decoder. The number of linear layers n in both the encoder and decoder is given by

$$n = \max(1, \lfloor \log_2(d_x) \rfloor).$$

The layer sizes of the encoder are created by always reducing the tensor size by t_e in each linear transformation $h_d - 1$ times, where t_e is calculated as

$$t_e = \min(-1, -\left\lfloor \frac{d_x - d_h}{n} \right\rfloor).$$

For the decoder, the input dimension d_x in the term is replaced by the output dimension d_y :

$$t_d = \min(-1, -\left\lfloor \frac{d_y - d_h}{n} \right\rfloor).$$

The decoder layer sizes are then calculated by reducing the layer size by t_d from the output layer in each linear transformation $h_d - 1$ times, but inverting the resulting array of layer sizes. This gives us a natural, symmetrical autoencoder structure for each test case (such as Figure 3.3), which is also used for following experiments. The results do not clearly indicate which autoencoder architecture is the best. Therefore we build upon prior research on mixed-type classification with attention [18, 13], where a feed-forward network that includes categorical features as well as continuous features is used. This means we choose the autoencoder that includes both categorical and continuous features.

We then compare our proposed architectures, an autoencoder with attention, the Transformer with an autoencoder and the FT-Transformer with an

	AE with categorical embeddings	AE with Attention	Transformer N=6	FT-Transformer N=6
Abalone	0.1503	0.1575	0.1345	0.1609
Auction Verification	0.8331	0.6422	0.6393	0.6290
Bank Marketing	0.7416	0.7044	0.5636	0.8796
Breast Cancer	0.9136	0.7277	0.8799	0.8184
Census Income	0.6740	0.7146	0.6812	0.7246
Credit Approval	0.5360	0.7335	0.7887	0.6141
Heart Disease	0.4448	0.3411	0.4515	0.3813
Soybean Disease	0.4751	0.2473	0.3114	0.2420

Figure 5.7: Comparison of accuracy between an autoencoder with categorical embeddings, an autoencoder using the attention mechanism, a 6-layer Transformer with an autoencoder and a 6-layer FT-Transformer with an autoencoder on clustering mixed-type datasets using k-means.

autoencoder. We choose $N = 6$ Transformer layers for both Transformers to utilize the transformer block heavily, but not introduce too much computational complexity. The autoencoder with categorical embeddings uses variable sized embeddings for each categorical feature (that are being concatenated lengthwise), where the embedding dimension is calculated based on the number of classes in that feature, as explained in Section 4.2. The attention autoencoder as well as both Transformer autoencoders use fixed size embeddings of size 32, which allows the embeddings to be stacked on top of each other, as explained in Section 4.3. The results (Figure 5.6 and 5.7) show that the attention can benefit clustering on some datasets, but leads to a loss in performance on other datasets.

At last, we use the models from the previous comparison, but use the deep clustering techniques DEC, IDEC and DCN (Section 3.3.3). While the accuracy results (Figure 5.9) are good, the Normalized Mutual Information results of these methods (Figure 5.8) do not indicate that deep clustering methods are particularly useful for clustering mixed-type tabular data.

		AE with categorical embeddings	Attention	Transformer N=6	FT-Transformer N=6
Abalone	DEC	0.1585	0.1630	0.1642	0.1520
	IDEC	0.1501	0.1623	0.1653	0.1660
	DCN	0.1574	0.1717	0.1671	0.1704
Auction Verification	DEC	0.0179	0.0203	0.0446	0.0000
	IDEC	0.0179	0.0203	0.0444	0.0000
	DCN	0.0179	0.0203	0.0484	0.0000
Bank Marketing	DEC	0.0195	0.0023	0.0000	0.0015
	IDEC	0.0172	0.0023	0.0000	0.0012
	DCN	0.0195	0.0023	0.0000	0.0006
Breast Cancer	DEC	0.7284	0.2185	0.5991	0.2062
	IDEC	0.7284	0.2185	0.2244	0.2062
	DCN	0.7333	0.2185	0.2062	0.2062
Census Income	DEC	0.0915	0.0663	0.0689	0.1333
	IDEC	0.0911	0.0663	0.0689	0.1333
	DCN	0.0915	0.0663	0.0689	0.1333
Credit Approval	DEC	0.0027	0.0000	0.0036	0.0725
	IDEC	0.0024	0.0000	0.0036	0.0725
	DCN	0.0024	0.0000	0.0036	0.0725
Heart Disease	DEC	0.1273	0.1599	0.1494	0.2434
	IDEC	0.1273	0.1599	0.1503	0.2434
	DCN	0.1268	0.1599	0.1558	0.2473
Soybean Disease	DEC	0.3878	0.5168	0.2605	0.2605
	IDEC	0.3866	0.6532	0.2684	0.2684
	DCN	0.3930	0.6306	0.2605	0.2605

Figure 5.8: Comparison of Normalized Mutual Information between an autoencoder with categorical embeddings, an autoencoder with attention, a 6-layer Transformer with an autoencoder and a 6-layer FT-Transformer with an autoencoder on clustering mixed-type datasets. DEC, IDEC and DCN were used as clustering methods.

		AE with categorical embeddings	Attention	Transformer N=6	FT-Transformer N=6
Abalone	DEC	0.1621	0.1286	0.1549	0.1858
	IDEC	0.1472	0.1384	0.1559	0.1678
	DCN	0.1676	0.1585	0.1685	0.2107
Auction Verification	DEC	0.6676	0.6760	0.6848	0.7445
	IDEC	0.6676	0.6760	0.6843	0.7445
	DCN	0.6676	0.6760	0.6569	0.7445
Bank Marketing	DEC	0.5712	0.7010	0.7602	0.7936
	IDEC	0.6162	0.7010	0.7602	0.7950
	DCN	0.5712	0.7010	0.7602	0.8172
Breast Cancer	DEC	0.9458	0.7467	0.8990	0.7350
	IDEC	0.9458	0.7467	0.7745	0.7350
	DCN	0.9473	0.7467	0.7350	0.7350
Census Income	DEC	0.6488	0.7414	0.6284	0.7236
	IDEC	0.6486	0.7414	0.6284	0.7236
	DCN	0.6488	0.7414	0.6284	0.7236
Credit Approval	DEC	0.5375	0.5299	0.5467	0.5697
	IDEC	0.5360	0.5299	0.5467	0.5697
	DCN	0.5360	0.5299	0.5467	0.5697
Heart Disease	DEC	0.2977	0.3445	0.3378	0.5652
	IDEC	0.2977	0.3445	0.3344	0.5652
	DCN	0.3010	0.3445	0.3445	0.5585
Soybean Disease	DEC	0.2544	0.4431	0.2278	0.2278
	IDEC	0.2544	0.5552	0.2295	0.2295
	DCN	0.2669	0.5125	0.2278	0.2278

Figure 5.9: Comparison of accuracy between an autoencoder with categorical embeddings, an autoencoder with attention, a 6-layer Transformer with an autoencoder and a 6-layer FT-Transformer with an autoencoder on clustering mixed-type datasets. DEC, IDEC and DCN were used as clustering methods.

	1st	2nd	3rd
Abalone	FT-Transformer (0.1750)	k-means one-hot (0.1740)	k-means (0.1718)
Auction Verification	AE with Column Embeddings (0.1234)	DCN Transformer (0.0484)	DEC Transformer(0.0446)
Bank Marketing	AE with Attention (0.0394)	k-means one-hot (0.0261)	k-means (0.0198)
Breast Cancer	Basic AE (0.7602)	k-means (0.7468)	k-means one-hot (0.7363)
Census Income	k-means one-hot (0.1850)	AE with Column Embeddings (0.1475)	FT-Transformer (0.1423)
Credit Approval	k-means (0.3131)	Transformer (0.3061)	AE with Attention (0.1833)
Heart Disease	DCN FT-Transformer (0.2473)	DEC/IDEC FT-Transformer (0.2434)	k-means (0.2046)
Soybean Disease	k-means one-hot (0.7102)	k-means (0.6722)	Gower distance (0.6695)

Figure 5.10: Best Normalized Mutual Information scores (in parenthesis) across all clustering methods. Classical methods are colored in green, methods using an autoencoder with k-means are colored cyan, methods using an autoencoder and attention or a transformer mechanism are colored yellow and methods using DEC, IDEC or DCN are colored magenta.

Chapter 6

Conclusion

Bibliography

- [1] Amir Ahmad and Shehroz S. Khan. Survey of state-of-the-art mixed data clustering algorithms. *IEEE Access*, 7:31883–31902, 2019.
- [2] Sercan O. Arik and Tomas Pfister. Tabnet: Attentive interpretable tabular learning, 2020.
- [3] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization, 2016.
- [4] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate, 2016.
- [5] Dana H. Ballard. Modular learning in neural networks. In *AAAI Conference on Artificial Intelligence*, 1987.
- [6] Christopher M. Bishop. *Neural Networks for Pattern Recognition*. Oxford University Press, Inc., USA, 1995.
- [7] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag, Berlin, Heidelberg, 2006.
- [8] Peter Bloem. Transformers from scratch. <https://peterbloem.nl/blog/transformers>, August 2018. Accessed: 2023-08-24.
- [9] Vadim Borisov, Tobias Leemann, Kathrin Sessler, Johannes Haug, Martin Pawelczyk, and Gjergji Kasneci. Deep neural networks and tabular data: A survey. *IEEE Transactions on Neural Networks and Learning Systems*, pages 1–21, 2022.
- [10] Zhengfan Chen, Shifei Ding, and Hou Haiwei. A novel self-attention deep subspace clustering. *International Journal of Machine Learning and Cybernetics*, 12, 08 2021.

- [11] Kamran Ghasedi Dizaji, Amirhossein Herandi, Cheng Deng, Weidong Cai, and Heng Huang. Deep clustering via joint convolutional autoencoder embedding and relative entropy minimization. In *2017 IEEE International Conference on Computer Vision (ICCV)*, pages 5747–5756, 2017.
- [12] Shiv Ram Dubey, Satish Kumar Singh, and Bidyut Baran Chaudhuri. Activation functions in deep learning: A comprehensive survey and benchmark. *Neurocomputing*, 503:92–108, 2022.
- [13] Yury Gorishniy, Ivan Rubachev, Valentin Khrulkov, and Artem Babenko. Revisiting deep learning models for tabular data. In *NeurIPS*, 2021.
- [14] J. C. Gower. A general coefficient of similarity and some of its properties. *Biometrics*, 27(4):857–871, 1971.
- [15] Xifeng Guo, Long Gao, Xinwang Liu, and Jianping Yin. Improved deep embedded clustering with local structure preservation. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17*, pages 1753–1759, 2017.
- [16] Xifeng Guo, En Zhu, Xinwang Liu, and Jianping Yin. Deep embedded clustering with data augmentation. In Jun Zhu and Ichiro Takeuchi, editors, *Proceedings of The 10th Asian Conference on Machine Learning*, volume 95 of *Proceedings of Machine Learning Research*, pages 550–565. PMLR, 14–16 Nov 2018.
- [17] Haiwei Hou, Shifei Ding, Xiao Xu, and Ling Ding. A novel clustering algorithm based on multi-layer features and graph attention networks. *Soft Computing*, 27:5553–5566, 2023.
- [18] Xin Huang, Ashish Khetan, Milan Cvitkovic, and Zohar Karnin. Tab-transformer: Tabular data modeling using contextual embeddings, 2020.
- [19] Zhexue Huang. Extensions to the k-means algorithm for clustering large data sets with categorical values. *Data Mining and Knowledge Discovery*, 2:283–304, 1998.
- [20] Anil K. Jain and Richard C. Dubes. *Algorithms for Clustering Data*. Prentice-Hall, Inc., USA, 1988.
- [21] Andras Janosi, William Steinbrunn, Matthias Pfisterer, and Robert Detrano. Heart Disease. UCI Machine Learning Repository, 1988. DOI: <https://doi.org/10.24432/C52P4X>.

- [22] Ron Kohavi. Census Income. UCI Machine Learning Repository, 1996. DOI: <https://doi.org/10.24432/C5GP7S>.
- [23] Jannik Kossen, Neil Band, Clare Lyle, Aidan N. Gomez, Tom Rainforth, and Yarin Gal. Self-attention between datapoints: Going beyond individual input-output pairs in deep learning. In *Neural Information Processing Systems*, 2021.
- [24] Bai Li. Clustering autoencoders: Comparing dec and dcn. <https://luckytoilet.wordpress.com/2019/07/07/clustering-autoencoders-comparing-dec-and-dcn/>, July 2019. Accessed: 2023-09-03.
- [25] Huifa Liao, Jie Hu, Tianrui Li, Shengdong Du, and Bo Peng. Deep linear graph attention model for attributed graph clustering. *Knowledge-Based Systems*, 246:108665, 2022.
- [26] S. Lloyd. Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28(2):129–137, 1982.
- [27] J. B. MacQueen. Some methods for classification and analysis of multivariate observations. In L. M. Le Cam and J. Neyman, editors, *Proc. of the fifth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 281–297. University of California Press, 1967.
- [28] Meena Mahajan, Prajakta Nimbhorkar, and Kasturi Varadarajan. The planar k-means problem is np-hard. *Theoretical Computer Science*, 442:13–21, 2012. Special Issue on the Workshop on Algorithms and Computation (WALCOM 2009).
- [29] Kolby Nottingham Markelle Kelly, Rachel Longjohn. The UCI Machine Learning Repository. <https://archive.ics.uci.edu>. Accessed: 2023-07-16.
- [30] Warren S McCulloch and Walter Pitts. A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5(4):115–133, 1943.
- [31] S. Moro, P. Rita, and P. Cortez. Bank Marketing. UCI Machine Learning Repository, 2012. DOI: <https://doi.org/10.24432/C5K306>.
- [32] Warwick Nash, Tracy Sellers, Simon Talbot, Andrew Cawthorn, and Wes Ford. Abalone. UCI Machine Learning Repository, 1995. DOI: <https://doi.org/10.24432/C55C7W>.

- [33] Elaheh Ordoni, Jakob Bach, Ann-Katrin Fleck, and Jakob Bach. Auction Verification. UCI Machine Learning Repository, 2022. DOI: <https://doi.org/10.24432/C52K6N>.
- [34] Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito, Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. Automatic differentiation in pytorch. In *NIPS-W*, 2017.
- [35] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [36] Zhihao Peng, Hui Liu, Yuheng Jia, and Junhui Hou. Deep attention-guided graph clustering with dual self-supervision. *IEEE Transactions on Circuits and Systems for Video Technology*, pages 1–1, 2022.
- [37] G. Philip and B. S. Ottaway. Mixed data cluster analysis: An illustration using cypriot hooked-tang weapons. *Archaeometry*, 25(2):119–133, 1983.
- [38] Quinlan Quinlan. Credit Approval. UCI Machine Learning Repository. DOI: <https://doi.org/10.24432/C5FS30>.
- [39] F. Rosenblatt. The perceptron - a perceiving and recognizing automaton. Technical Report 85-460-1, Cornell Aeronautical Laboratory, Ithaca, New York, January 1957.
- [40] D. E. Rumelhart, G. E. Hinton, and R. J. Williams. *Learning Internal Representations by Error Propagation*, page 318–362. MIT Press, Cambridge, MA, USA, 1986.
- [41] Gowthami Somepalli, Micah Goldblum, Avi Schwarzschild, C. Bayan Bruss, and Tom Goldstein. Saint: Improved neural networks for tabular data via row attention and contrastive pre-training, 2021.
- [42] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, NIPS’17, page 6000–6010, Red Hook, NY, USA, 2017. Curran Associates Inc.
- [43] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, Matt Haberland, Tyler Reddy, David Cournapeau, Evgeni Burovski, Pearu Peterson, Warren Weckesser, Jonathan Bright, Stéfan J. van der Walt, Matthew Brett,

- Joshua Wilson, K. Jarrod Millman, Nikolay Mayorov, Andrew R. J. Nelson, Eric Jones, Robert Kern, Eric Larson, C J Carey, İlhan Polat, Yu Feng, Eric W. Moore, Jake VanderPlas, Denis Laxalde, Josef Perktold, Robert Cimrman, Ian Henriksen, E. A. Quintero, Charles R. Harris, Anne M. Archibald, Antônio H. Ribeiro, Fabian Pedregosa, Paul van Mulbregt, and SciPy 1.0 Contributors. SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods*, 17:261–272, 2020.
- [44] William Wolberg. Breast Cancer Wisconsin (Original). UCI Machine Learning Repository, 1992. DOI: <https://doi.org/10.24432/C5HP4Z>.
- [45] Junyuan Xie, Ross Girshick, and Ali Farhadi. Unsupervised deep embedding for clustering analysis, 2016.
- [46] Bo Yang, Xiao Fu, Nicholas D. Sidiropoulos, and Mingyi Hong. Towards k-means-friendly spaces: Simultaneous deep learning and clustering, 2017.
- [47] Sheng Zhou, Hongjia Xu, Zhuonan Zheng, Jiawei Chen, Zhao li, Jiajun Bu, Jia Wu, Xin Wang, Wenwu Zhu, and Martin Ester. A comprehensive survey on deep clustering: Taxonomy, challenges, and future directions, 2022.